

Defense Brief — Representation Convergence Project

Companion to the 61-page report. One page of framing per topic; every number cited is measured and appears in the report with its generator notebook.

1 · Assumptions (what the results are conditional on)

Item-level correspondence exists and is known for supervised fits. Row i of every array is the same passage or image in every model. This is constructed (nested sorted-id prefixes, deterministic caption recipe) and verified numerically wherever two caches meet; C.12 is the one experiment that deliberately removes it.

Linear maps, closed form. Every fitted map is ridge regression solved in closed form - no gradients, no iteration - so 'related by a linear transformation' carries no asterisk. The MLP contrast (F2) is the deliberate exception and converged to the linear answer exactly.

Frozen encoders. No pretrained weight is modified anywhere in the evidence chain; F3, which would unfreeze one, is built but gated shut to preserve that statement.

Held-out discipline includes the transforms. Splits precede fitting; whitening covariances, hub bases, and alpha choices come from training rows only. Gallery size is 1,000 everywhere, so chance is 0.001 and R@K columns are comparable across experiments.

Geometry diagnostics are trusted as measurements. Effective rank (spectral entropy) and mean pair-cosine are treated as valid summaries of a space's shape; C.14's perfect rank ordering is itself evidence the trust is earned.

2 · Restrictions (scope the claims do not exceed)

claim	measured scope	explicitly NOT claimed
convergence is linear	GPT-2/Pythia; MobileCLIP/SigLIP; DINOv2 x3; 4 text encoders	all architectures, all scales
alignment scales with capacity	DINOv2 family, 22M-304M (13.8x), compressive at the top	behaviour beyond 304M; other families
isotropy rescue	GPT-2 target, two image scales (102.6%/101.4%)	sufficiency - C.14 bounds it
unsupervised recovery	39.7% exact @ 595x chance; 40% of ceiling	a deployable pairing-free system
shared hub	5 encoders, 2 families, COCO, English; transfer 93-96%, cross-lineage and cross-architecture	a universal space for all models
deployment numbers	COCO val/train2017 domain; 1,000-item galleries	open-world scale or other domains

One language, one image domain, one caption corpus throughout. Every 'per cent of ceiling' is against a ceiling re-measured under the same protocol and caption treatment - never a borrowed number.

3 · Tuning (every knob, and how it was set)

knob	how set	guard against overfitting the knob
ridge alpha	validation sweep 1e-3..10 per experiment	chosen on train-internal validation, never on eval
pooling	fixed cls+patch across the scale sweep	the one mixed run (base CLS) excluded from comparisons
N (pairs)	scaled with encoder width -> 11.1 rows/dim at every sweep point	starvation biases DOWN, so it cannot fake the trend
GW epsilon	swept per run; optimum MOVES with N (5e-3@800 -> 1e-3@1500)	reported as a sweep, never a single tuned value
HUB_DIM x alpha	2-D sweep, transfer-first, rows/dim ≥ 5 enforced	pre-registered thresholds; cliff reported as finding
gates	90%-of-ceiling; 85% transitivity; 50% transfer; ≥ 5 rows/dim	registered before results; 89.0% reported as NOT met

4 · Defense — the hard questions, with answers

“Isn't whitening cheating?” It is label-free (uses no pairs, no labels), fitted on training targets only, and linear - so it is an instance of the thesis, not an exception. The full map is whitened composed with ridge: still one matrix.

“Why is the whitened R-squared 0.014?” Whitening destroys magnitudes and perfects angles. R-squared asks about magnitudes; retrieval asks about angles. One system, two metrics, opposite verdicts - which is the project's metric-geometry lesson in its purest form.

“You said isotropy helps - then why does it destroy C.12?” Reading versus finding. Cosine reads a KNOWN correspondence: uneven variance is noise. Gromov-Wasserstein FINDS an unknown one from distance structure: uneven variance is the signal that makes points distinctive. Same property, opposite prescriptions - both measured, twice each.

“Doesn't C.14's residual spread undermine C.11?” It bounds it. C.11's mechanism predicted two unseen encoders' ordering perfectly (rank-corr +1.00). The residue after whitening (0.110) measures what contrastive training adds BEYOND the frame: the arrangement of meaning across directions. Necessary, not sufficient - a stronger theory than the original.

“Why did the hub break between 512 and 768 dimensions?” I don't know - and I tested SIX answers to destruction. Transfer 0.489 -> 0.312, a -0.177 step against +0.015 before it. Local-structure loss: FALSIFIED (-0.018 / -0.030, six times too small, and smooth). Spectral amplification, my own account: FALSIFIED ($1/\sqrt{\lambda}$ decelerates, 1.45x -> 1.26x). Four further statistics: all FALSIFIED (D 0.62-0.80 vs transfer's 4.03). Not power (rows/dim ≥ 5 both sides). Unexplained, not argued away.

“Is the transfer result circular - the head just learned the task?” The random-map control replaces only the hub map and collapses to chance (0.001), so the head alone carries nothing. And absolute R@1 RISES on unseen encoders, which a task-memorising head cannot produce.

“Have the artifacts themselves been checked?” Yes, in one automated run after all experiments were complete: 47 checks across every stored artifact - statistical power, row alignment by shuffle test, sanity, target geometry, and cache-key integrity. NO hard failures. Every fit sits between 5.9 and 9.8 rows per input dimension; shuffle gaps run 0.545 to 0.863 against a 0.2 threshold; all three encoder caches match the configuration named in their filenames. The only two flags are the GPT-2 caption space, whose collapse is a documented finding rather than a defect - and the battery reproduced its signature independently in two separately generated files. (Appendix D.5.)

“How do I know the numbers are right?” Five documented self-corrections via measurement (starvation rerun, synthetic-figure retraction, matched ceiling, wrong reference arm twice - now derived from filenames, whitening-for-GW reversal), plus width-tagged checkpoints, transfer checks against known pairs, and AST-validated notebooks. The error CLASS is real; the defenses are structural.


5 · Achievements — the nine findings

#	finding	headline evidence	section
1	convergence is linear	MLP = ridge exactly at 74 rows/dim (0.9052 = 0.9052)	C.9/F
2	related, not rigid	Procrustes 0.038 vs ridge 0.592; 69x singular-value spread	4, B
3	deployable	95% of ceiling @ 0.79 MB; verif AUC 0.999 / rank 0.987 / membership 0.889	4, D, D.2
4	cross-modal, no joint training	40.9% of matched ceiling, 358x chance	C.8-C.9
5	alignment scales with capacity	+15.3 pts of ceiling per decade, 6 monotone measures, compressive	C.10
6	isotropy, not similarity training	whitening -> 102.6% / 101.4% of bge at two scales; rho +0.94	C.11
7	correspondence is self-identifying	GW 39.7% exact = 595x chance, 0 pairs; + 2 measured negatives	C.12
8	one shared space, used	hub transfer 95.9%/92.9% of native on unseen encoders; control at chance	C.13
9	the mechanism predicts	pre-registered ladder: rank-corr +1.00; isotropy necessary-not-sufficient	C.14

Plus the recurring meta-finding, measured three times: isotropy helps reading a correspondence and hurts finding or losslessly carrying one (C.11 vs C.12 vs the C.13 width cliff).


6 · Flow: how E1.3 closes the loop

E1.3 — the encoder ladder: one image side, four text objectives, predictions registered first



Flow: the hub, built and demonstrated

G1/G2 — five encoders, one 512-d hub: the shared space built, measured, and demonstrated live



7 · Conclusion

The thesis - independently trained models share representations up to a linear transformation - was tested in every posture available: pairwise and supervised (it holds, and an MLP adds exactly nothing); across modalities without joint training (it holds at 41-58 per cent of a jointly-trained ceiling, rising with capacity at a measured slope); with the metric geometry itself interrogated (apparent failures were coordinate-frame artifacts, removable by a free linear transform, with the residue bounded by a pre-registered prediction test); with the pairing itself withheld (the correspondence partially identifies itself from shape alone); and constructively (one 512-dimensional linear hub serves five encoders, and a component trained on one transfers to encoders it never saw at 93-96 per cent of native performance).

Two things distinguish the project beyond the results. First, its negatives are load-bearing and mechanistic: ICP converging on its own errors, whitening destroying discovery, the hub-width cliff - each measured, each explaining the boundary it marks. Second, its errors are part of the record: five self-corrections reached by measurement rather than argument, which is the difference between a demonstration and research. The final state of the claim is narrower and stronger than the initial one: the shared structure is real, linear, discoverable, and usable - and the coordinate frame that makes it readable is cheap, while the semantic arrangement that makes it meaningful is what training actually buys.